



**Budapest University of Technology and Economics**  
Faculty of Electrical Engineering and Informatics  
Department of Control Engineering and Information Technology

# Algorithms for Real-Time Physics Simulation

BACHELOR'S THESIS

*Author*  
Márton Biró

*Advisor*  
dr. Márton Vaitkus

May 24, 2026

# Contents

|                                                             |           |
|-------------------------------------------------------------|-----------|
| <b>Kivonat</b>                                              | <b>i</b>  |
| <b>Abstract</b>                                             | <b>ii</b> |
| <b>1 Introduction</b>                                       | <b>1</b>  |
| 1.1 Motivation . . . . .                                    | 1         |
| 1.2 Contribution . . . . .                                  | 2         |
| 1.3 Hypotheses . . . . .                                    | 3         |
| 1.4 Evaluation axes . . . . .                               | 3         |
| 1.5 Thesis structure . . . . .                              | 3         |
| <b>2 Background and theoretical foundations</b>             | <b>5</b>  |
| 2.1 Time-stepping schemes for physical simulation . . . . . | 5         |
| 2.1.1 Explicit Euler and why it fails . . . . .             | 6         |
| 2.1.2 Implicit Euler . . . . .                              | 6         |
| 2.2 Implicit Euler as energy minimisation . . . . .         | 6         |
| 2.3 Mass-spring systems . . . . .                           | 7         |
| 2.3.1 Distance, stretch, and bending springs . . . . .      | 7         |
| 2.3.2 Limitations of mass-spring . . . . .                  | 7         |
| 2.4 Solver structure: Gauss–Seidel vs. Jacobi . . . . .     | 7         |
| 2.5 The primal/dual distinction . . . . .                   | 8         |
| <b>3 Related work</b>                                       | <b>9</b>  |
| 3.1 From PBD to XPBD . . . . .                              | 9         |
| 3.2 Projective dynamics and descent methods . . . . .       | 10        |
| 3.3 Vertex Block Descent (VBD) and successors . . . . .     | 10        |
| 3.4 Prior comparison: Ma et al. (2024) . . . . .            | 10        |
| 3.5 Reference materials . . . . .                           | 11        |
| <b>4 The two algorithms</b>                                 | <b>12</b> |

|          |                                                                     |           |
|----------|---------------------------------------------------------------------|-----------|
| 4.1      | XPBD: dual-variable position projection . . . . .                   | 12        |
| 4.1.1    | Derivation . . . . .                                                | 12        |
| 4.1.2    | Algorithm . . . . .                                                 | 12        |
| 4.1.3    | Characteristics . . . . .                                           | 13        |
| 4.2      | VBD: primal block-coordinate descent . . . . .                      | 13        |
| 4.2.1    | Derivation . . . . .                                                | 13        |
| 4.2.2    | Algorithm . . . . .                                                 | 13        |
| 4.2.3    | Characteristics . . . . .                                           | 14        |
| 4.3      | Side-by-side: where the algorithms differ . . . . .                 | 14        |
| <b>5</b> | <b>Methodology and implementation</b>                               | <b>15</b> |
| 5.1      | Platform and environment . . . . .                                  | 15        |
| 5.2      | Integration architecture . . . . .                                  | 15        |
| 5.3      | How the comparison is kept fair . . . . .                           | 16        |
| 5.3.1    | Stiffness $\leftrightarrow$ compliance mapping . . . . .            | 16        |
| 5.3.2    | Matched compute budget — and the substep/iteration caveat . . . . . | 16        |
| 5.3.3    | Matched gravity, $dt$ , and damping . . . . .                       | 17        |
| 5.3.4    | Disabling VBD’s accelerations for the core comparison . . . . .     | 17        |
| 5.4      | Damping models . . . . .                                            | 17        |
| 5.5      | Self-collision handling . . . . .                                   | 18        |
| 5.6      | Parallelisation (forward-looking) . . . . .                         | 18        |
| <b>6</b> | <b>Experiments and results</b>                                      | <b>19</b> |
| 6.1      | Computational efficiency . . . . .                                  | 19        |
| 6.1.1    | Wall-clock . . . . .                                                | 19        |
| 6.1.2    | Frame time vs. node count . . . . .                                 | 20        |
| 6.2      | Robustness and stability . . . . .                                  | 20        |
| 6.2.1    | High mass ratio at the end of a pendulum . . . . .                  | 20        |
| 6.2.2    | High mass ratio chain . . . . .                                     | 20        |
| 6.2.3    | High stiffness ratio chain . . . . .                                | 20        |
| 6.2.4    | Self-collision: cloth folded onto itself . . . . .                  | 20        |
| 6.3      | Physical realism . . . . .                                          | 21        |
| 6.3.1    | Newton reference . . . . .                                          | 21        |
| 6.3.2    | Harmonic oscillator . . . . .                                       | 21        |
| 6.3.3    | Energy tracking . . . . .                                           | 21        |
| 6.3.4    | Hanging cloth: stretched area vs. reference . . . . .               | 21        |
| 6.3.5    | Damping cross-check . . . . .                                       | 22        |

|          |                                                                        |           |
|----------|------------------------------------------------------------------------|-----------|
| 6.4      | Flexibility . . . . .                                                  | 22        |
| 6.4.1    | Iteration-count independence of stiffness . . . . .                    | 22        |
| 6.4.2    | Stiffness-ratio test . . . . .                                         | 22        |
| 6.4.3    | Ease of adding a new constraint or force . . . . .                     | 22        |
| <b>7</b> | <b>Discussion</b>                                                      | <b>23</b> |
| 7.1      | Synthesis across axes . . . . .                                        | 23        |
| 7.2      | Verdicts on H1–H3 . . . . .                                            | 23        |
| 7.3      | Threats to validity . . . . .                                          | 24        |
| 7.3.1    | Single-threaded execution . . . . .                                    | 24        |
| 7.3.2    | Mass-spring discretisation only . . . . .                              | 24        |
| 7.3.3    | Gauss–Seidel order dependence . . . . .                                | 24        |
| 7.3.4    | Mismatched collision models . . . . .                                  | 25        |
| 7.3.5    | Fixed $dt = 1/30$ s . . . . .                                          | 25        |
| 7.4      | Limitations . . . . .                                                  | 25        |
| <b>8</b> | <b>Conclusion and future work</b>                                      | <b>26</b> |
| 8.1      | Summary . . . . .                                                      | 26        |
| 8.2      | Future work . . . . .                                                  | 26        |
| 8.2.1    | Multithreaded and GPU comparison . . . . .                             | 26        |
| 8.2.2    | Extend research to other dynamic types . . . . .                       | 26        |
| 8.2.3    | FEM constraints . . . . .                                              | 27        |
| 8.2.4    | AVBD and post-2024 variants . . . . .                                  | 27        |
|          | <b>Acknowledgements</b>                                                | <b>28</b> |
|          | <b>Bibliography</b>                                                    | <b>29</b> |
|          | <b>Appendix</b>                                                        | <b>31</b> |
| A.1      | Declaration on the Use of Generative Artificial Intelligence . . . . . | 32        |

## STUDENT DECLARATION

I, *Márton Biró*, the undersigned, hereby declare that the present Bachelor's Thesis work has been prepared by myself and without any unauthorized help or assistance. Only the specified sources (references, tools, etc.) were used. All parts taken from other sources word by word, or after rephrasing but with identical meaning, were unambiguously identified with explicit reference to the sources utilized. In the Appendix, I have clearly indicated whether I used artificial intelligence tools in preparing the thesis; if so, I have specified the manner and extent of their use in the table. I acknowledge that I bear full responsibility for any content generated by artificial intelligence – regardless of the extent of its use.

I agree that the basic information of my present work (author(s), title, abstracts in English and in a second language, year of preparation, name of advisor(s)) will be made publicly available on the internet by BME VIK, and the full text of the work will be available to persons registered in the BME Directory. I declare that the submitted hardcopy of the thesis work and its electronic version are identical.

Full text of thesis works classified upon the decision of the Dean will be published after a period of three years.

Budapest, May 24, 2026

---

*Márton Biró*  
student

# Kivonat

\*TODO\*

# Abstract

\*TODO\*

# Chapter 1

## Introduction

### 1.1 Motivation

- Real-time simulation of deformable objects (cloth, soft bodies, hair) is a baseline expectation in games, VR, and AR.
- The simulator must deliver believable dynamics within a tight frame budget: at 60 fps a frame allows  $\sim 16$  ms for all physics, rendering, and gameplay logic, with physics getting only a fraction.
- I focus on deformable bodies rather than rigid bodies because:
  - the two studied algorithms differ most visibly on deformable material — rigid-body simulation hides the per-constraint convergence behaviour I want to measure;
  - deformable simulation is the strictly harder case (many more degrees of freedom and constraints per frame), stressing both solvers so their failure modes become legible.

#### The accuracy/efficiency trade-off.

- Every physics-based simulator sits on a line between *physical accuracy* and *computational efficiency*; gaining one costs the other.
- Table 1.1 sketches this spectrum, ordered from most accurate to most efficient.
- XPBD and VBD both sit near the efficient end — where real-time applications live — and the thesis asks which makes the better trade-off, and under what conditions.

Designing such a simulator is hard for three recurring reasons:

- *Stiffness*. Cloth and soft tissue have very stiff constraints (almost-inextensible fabric, near-rigid bones); stiff systems make explicit integrators unstable unless time steps become impractically small.
- *Stability*. Even implicit integrators can blow up under poorly conditioned constraints, large mass ratios between connected particles, or contact-induced velocity discontinuities.



| Method                                  | Accuracy             | Cost         |
|-----------------------------------------|----------------------|--------------|
| Newton’s method (implicit Euler)        | reference            | offline only |
| Projective Dynamics                     | high                 | expensive    |
| Vertex Block Descent (VBD)              | moderate–high        | real-time    |
| Extended Position-Based Dynamics (XPBD) | moderate             | real-time    |
| Position-Based Dynamics (PBD)           | low (stiffness-dep.) | cheap        |
| Explicit Euler                          | poor (unstable)      | cheapest     |

**Table 1.1:** Deformable-simulation methods ordered from most accurate to most efficient. XPBD and VBD — the subject of this thesis — both target the real-time regime.

- *Strict time constraints.* Methods that converge well offline (Newton’s method, projective dynamics) are too expensive per frame; cheap real-time methods often trade convergence quality for predictable cost.

Within the real-time band:

- XPBD [6] has been a workhorse for years.
- VBD [1] is a recent, promising addition not yet evaluated head-to-head against XPBD.
- The two occupy opposite sides of the *primal/dual* duality identified by Macklin et al. [5]: XPBD solves for Lagrange multipliers (dual variables on constraints), VBD solves directly for vertex positions (primal variables).
- No controlled comparison between them currently exists.

## 1.2 Contribution

- The central contribution is, to the best of my knowledge, the *first faithful XPBD-vs-VBD comparison*:
  - both algorithms implemented in the same Unity-based codebase,
  - applied to the same mass-spring material model,
  - run under a matched compute budget,
  - evaluated against a high-iteration Newton solve serving as numerical ground-truth reference.
- This sharpens the closest prior comparison, Ma et al. [4], which contrasts *plain* PBD (not XPBD) with a non-canonical VBD implementation and reports only frame-rate scaling.

I distinguish this work by:

1. using XPBD, the constraint-stiffness-independent variant any practitioner would deploy today;

2. implementing the canonical VBD formulation, including its block-coordinate-descent interpretation;
3. matching compute budgets and material parameters rather than iteration counts (one XPBD substep is not directly comparable to one VBD iteration);
4. adding a Newton ground truth, enabling reporting of *solution error* rather than only relative performance.

### 1.3 Hypotheses

- The primal/dual theory in Chapter 2 predicts XPBD and VBD fail in characteristically different regimes.
  - Three testable hypotheses are stated up front; the experiments in Chapter 6 are designed to confirm or refute them.
- H1.** *XPBD degrades under high mass ratios.* Because XPBD solves dual variables one constraint at a time, large mass disparities along a constraint chain amplify the local update error and propagate slowly.
- H2.** *VBD degrades under high stiffness ratios.* VBD’s primal block-coordinate update relies on a locally well-conditioned Hessian; a wide range of stiffnesses meeting at a single vertex produces an ill-conditioned block and slows convergence.
- H3.** *VBD’s per-iteration convergence is more stiffness-robust.* Because the primal update sees the true elastic Hessian rather than a linearised constraint, raising material stiffness should not slow VBD as much as it slows XPBD.

### 1.4 Evaluation axes

I evaluate both solvers along four axes that together capture the trade-offs a practitioner cares about:

- Efficiency.** Wall-clock cost per frame as a function of mesh size, substep/iteration count, and target quality.
- Robustness.** Behaviour under stress: large time steps, single iterations, extreme mass ratios, extreme stiffness ratios, self-contact.
- Physical realism.** Distance from a Newton reference, energy behaviour, damping behaviour, and agreement with analytic solutions on a harmonic oscillator and a draped cloth.
- Flexibility.** Ease of changing material stiffness, adding new constraint types, and reasoning about parameters.

### 1.5 Thesis structure

- Chapter 2: theoretical foundation — variational form of implicit Euler, mass-spring energies, the primal/dual lens that makes XPBD and VBD comparable.

- Chapter 3: reviews PBD, XPBD, VBD literature and positions this work against Ma et al. [4].
- Chapter 4: states the two algorithms in canonical form.
- Chapter 5: the Unity implementation and the design decisions that keep the comparison fair.
- Chapter 6: results along the four axes.
- Chapter 7: synthesis across axes, verdicts on H1–H3, threats to validity.
- Chapter 8: closing with future-work directions.

## Chapter 2

# Background and theoretical foundations

- This chapter establishes the mathematical scaffolding for the rest of the thesis.
- The central object is the variational form of implicit Euler, which unifies XPBD and VBD as two approximate minimisers of the same energy functional, and therefore justifies a shared Newton reference as ground truth.

### 2.1 Time-stepping schemes for physical simulation

- Real-time physics reduces to advancing a state  $(\mathbf{x}, \mathbf{v})$  — positions and velocities of  $n$  particles — forward in time.
- Three families of approaches have been used in computer graphics:

#### **Newtonian, force-based.**

- Forces  $\mathbf{f}(\mathbf{x}, \mathbf{v}, t)$  produce accelerations via  $\ddot{\mathbf{x}} = \mathbf{M}^{-1}\mathbf{f}$ , integrated by some scheme.

#### **Impulse-based.**

- Velocities are manipulated directly via collision or constraint impulses; popular for rigid bodies.
- Not pursued here: awkward for distributed deformation, and the primal/dual framework of Section 2.5 is the right lens for comparing XPBD and VBD.

#### **Position-based.**

- Positions are manipulated directly; velocities are recovered as finite differences. This is the lineage XPBD belongs to.

### 2.1.1 Explicit Euler and why it fails

The explicit (forward) Euler scheme is

$$\mathbf{v}_{t+h} = \mathbf{v}_t + h \mathbf{M}^{-1} \mathbf{f}(\mathbf{x}_t), \quad (2.1)$$

$$\mathbf{x}_{t+h} = \mathbf{x}_t + h \mathbf{v}_{t+h}. \quad (2.2)$$

- For a linear spring with stiffness  $k$  and mass  $m$ , stability requires roughly  $h < 2\sqrt{m/k}$ .
- For cloth, where  $k$  may exceed  $10^4$  N/m, this drives  $h$  down to microseconds — incompatible with a 16 ms real-time frame budget.
- Velocity Verlet and symplectic-Euler variants still retain a stability limit driven by the stiffest mode.

### 2.1.2 Implicit Euler

The (backward) implicit Euler scheme evaluates forces at the unknown end-of-step state:

$$\mathbf{v}_{t+h} = \mathbf{v}_t + h \mathbf{M}^{-1} \mathbf{f}(\mathbf{x}_{t+h}), \quad (2.3)$$

$$\mathbf{x}_{t+h} = \mathbf{x}_t + h \mathbf{v}_{t+h}. \quad (2.4)$$

- Unconditionally stable for linear systems, dissipative for non-linear ones.
- Cost: a non-linear root-finding problem each step.

## 2.2 Implicit Euler as energy minimisation

Eliminating the velocity equation gives the well-known variational form of implicit Euler: the new positions  $\mathbf{x}$  minimise

$$G(\mathbf{x}) = \frac{1}{2h^2} \|\mathbf{x} - \mathbf{y}\|_{\mathbf{M}}^2 + E(\mathbf{x}), \quad (2.5)$$

where

$$\mathbf{y} = \mathbf{x}_t + h \mathbf{v}_t + h^2 \mathbf{M}^{-1} \mathbf{f}_{\text{ext}}$$

- $\mathbf{y}$  is the *inertial target* (the position reached under external forces alone).
- $\|\cdot\|_{\mathbf{M}}$  is the mass-weighted norm;  $E(\mathbf{x})$  is the elastic potential energy.
- Equation (2.5) is central: **both XPBD and VBD are approximate minimisers of the same  $G(\mathbf{x})$ .**
- XPBD reaches a stationary point by descending on the Lagrange-dual problem; VBD reaches one by block-coordinate descent on the primal.
- This shared target licenses a Newton solve of (2.5) to serve as unambiguous ground truth in Chapter 6.

## 2.3 Mass-spring systems

- I use a mass-spring discretisation of cloth: vertices carry lumped masses, edges of the cloth mesh carry springs.
- For a spring between particles  $i$  and  $j$  with rest length  $\ell_0$  and stiffness  $k$ , the constraint function and energy are:

$$C(\mathbf{x}_i, \mathbf{x}_j) = \|\mathbf{x}_i - \mathbf{x}_j\| - \ell_0, \quad (2.6)$$

$$E(\mathbf{x}_i, \mathbf{x}_j) = \frac{1}{2} k C^2. \quad (2.7)$$

### 2.3.1 Distance, stretch, and bending springs

Two spring types, both expressed through the distance-spring energy above:

- *Stretch springs* between mesh-adjacent vertices oppose in-plane deformation.
- *Bending springs* are modelled as long-range distance springs between vertices opposite across an edge (the “diagonal” of the two-triangle quad).
- Note: this is *not* a dihedral-angle constraint — I deliberately stay within a single constraint form, since the goal is a controlled comparison, not a state-of-the-art cloth model.

### 2.3.2 Limitations of mass-spring

- Mass-spring is cheap and easy to differentiate, but is well known to be *mesh-dependent*: effective macroscopic stiffness depends on the discretisation.
- Therefore conclusions about absolute material parameters are not transferable to FEM-based models.
- Conclusions about *relative* solver behaviour — my subject — are transferable.

## 2.4 Solver structure: Gauss–Seidel vs. Jacobi

- Both XPBD’s constraint sweep and VBD’s vertex-block sweep are Gauss–Seidel-style iterations: each local update uses the most recent neighbour values, not start-of-iteration values.
- A Jacobi variant updates all blocks from the start-of-iteration state in parallel: slower convergence in serial, but maps better to the GPU.
- Consequence (revisited in Section 7.3): both solvers are **order-dependent**.
  - Constraint-traversal order (XPBD) and vertex-traversal order (VBD) both affect convergence rate and, in pathological cases, the converged solution.
  - My experiments fix a deterministic order for reproducibility.

## 2.5 The primal/dual distinction

- The unifying framing is the primal/dual perspective due to Macklin et al. [5].
- Minimising  $G(\mathbf{x})$  in (2.5) can be approached two ways:

### Primal.

- Directly minimise over positions  $\mathbf{x}$ .
- A vertex-by-vertex block-coordinate descent on  $G$  is exactly VBD.

### Dual.

- Reformulate the energy through compliant constraints (each spring is a soft constraint with compliance  $\alpha = 1/k$ ) and solve for Lagrange multipliers  $\boldsymbol{\lambda}$  on each constraint.
- XPBD is a Gauss–Seidel iteration on the resulting saddle-point system.

This duality predicts the hypotheses in Section 1.3:

- Dual variables couple constraints sharing a particle, so mass disparities at a shared particle hurt XPBD (H1).
- The primal Hessian block at a vertex collects all spring stiffnesses meeting there, so widely varying stiffnesses produce an ill-conditioned block and hurt VBD (H2).
- Stiffness-robust per-iteration convergence in VBD (H3) follows from the primal Hessian containing the true elastic operator, while the constraint linearisation underlying XPBD discards second-order information.

## Chapter 3

# Related work

- This chapter surveys the position-based and primal-descent lineages whose endpoints — XPBD and VBD — this thesis compares.
- It positions my work relative to the closest prior comparison.

### 3.1 From PBD to XPBD

- Position-Based Dynamics (PBD), introduced by Müller et al. [9], targets real-time cloth and soft-body simulation in games.
- Each iteration projects positions onto constraint manifolds (e.g. fixed edge lengths) via Gauss–Seidel sweeps, recovering velocities by finite difference.
- PBD is unconditionally stable but has a known flaw: *constraint stiffness depends on iteration count and time step*.
  - Practitioners cannot specify a Young’s modulus directly; they specify a stiffness coefficient that interacts unpredictably with solver settings.
- XPBD, introduced by Macklin et al. [6], fixes this by augmenting each constraint with a compliance term  $\alpha = 1/k$  and solving for the associated Lagrange multiplier  $\lambda$ .
- The resulting iteration solves an implicit-Euler problem with arbitrary elastic energies, and constraint stiffness becomes a true material parameter independent of iteration count.

#### Small steps in physics simulation.

- Macklin et al. [7] (the “Small Steps” note) later argued that XPBD converges most reliably with *many short substeps and a single iteration per substep*, rather than a single long step with many iterations.
- This is now a standard XPBD configuration and is the default I adopt in Chapter 5: it follows the recommendation and exposes XPBD’s per-substep cost as the primary performance lever.



## 3.2 Projective dynamics and descent methods

- Projective dynamics recasts implicit Euler as a quadratic minimisation alternating between a global linear solve and local constraint projections.
- Variants and accelerations (descent methods, Chebyshev acceleration, ADMM) bridge between projective dynamics and position-based methods.
- They provide the theoretical bridge to VBD: VBD can be read as a *block-coordinate descent* on the same variational energy (2.5) that projective dynamics minimises, but with single-vertex blocks rather than a global solve.

## 3.3 Vertex Block Descent (VBD) and successors

- Vertex Block Descent, due to Chen et al. [1], treats each vertex as a tiny optimisation block: at each iteration, every vertex moves to the minimiser of  $G$  with all other vertices held fixed.
- This requires the local  $3 \times 3$  Hessian of the energies sharing the vertex and a Newton step on that block.

Critical features of the original VBD paper:

- *Chebyshev acceleration*, an over-relaxation that significantly speeds convergence at the cost of energy guarantees.
- *Adaptive initialisation*, seeding the iteration from a blend of inertial and previous-step positions.
- *Graph-colouring parallelism*: vertices sharing no constraint can be updated independently, enabling massive parallelism.
- A follow-up, Augmented VBD (AVBD) [2], has improved stiffness-ratio handling and convergence behaviour.
- I restrict the comparison to the canonical 2024 VBD since AVBD post-dates my implementation window; I cite AVBD in Section 6.4 as the most relevant counter-evidence on the stiffness-ratio axis.

## 3.4 Prior comparison: Ma et al. (2024)

- The most directly comparable prior work is Ma et al. [4], who compare cloth simulation under PBD and VBD in Unity.
- Their study established that VBD scales better than PBD with node count, and identified maximum sustainable node counts at 30 fps.

My work differs in three substantial ways:

1. **XPBD, not PBD.** Plain PBD has the stiffness-coupling flaw that XPBD was designed to fix, so a PBD-vs-VBD comparison conflates algorithmic differences with this known bug. Using XPBD isolates the primal/dual distinction.
  2. **Canonical VBD.** Ma et al.’s pseudocode departs in details (initialisation, Chebyshev handling) from [1]. I follow canonical VBD and disable Chebyshev / adaptive init for the core comparison, so I read off algorithm properties, not acceleration tricks.
  3. **Beyond frame rate.** Ma et al. report fps and node counts. I add a Newton ground truth, an equal-wall-clock quality frontier, and targeted stress tests (mass ratio, stiffness ratio, large  $dt$ ) keyed to the primal/dual hypotheses.
- In short, this thesis answers a question Ma et al. explicitly left as future work: how does VBD compare to *XPBD*, and at *matched quality* rather than only at matched node counts?

### 3.5 Reference materials

- The Physics-Based Simulation book of Li et al. [3] provided canonical derivations and pedagogical notes; their treatment of implicit Euler as variational minimisation directly informs the framing in Section 2.2.
- The *Ten Minute Physics* tutorial series [8] was used as a sanity-check reference for XPBD implementation correctness.

# Chapter 4

## The two algorithms

- This chapter states XPBD and VBD in the form actually implemented in Chapter 5.
- Both target the same variational problem (2.5) but reach it through opposite descent directions.
- The contrast made explicit here drives the experimental design in Chapter 6.

### 4.1 XPBD: dual-variable position projection

#### 4.1.1 Derivation

- For each elastic spring, associate a constraint  $C(\mathbf{x}) = 0$  at zero deformation and a compliance  $\alpha = 1/k$ .
- The augmented Lagrangian of (2.5) with a single substep  $h$  has stationarity conditions:

$$\mathbf{M}(\mathbf{x} - \mathbf{y}) = h^2 \nabla C(\mathbf{x})^\top \boldsymbol{\lambda}, \quad (4.1)$$

$$C(\mathbf{x}) + \tilde{\alpha} \boldsymbol{\lambda} = \mathbf{0}, \quad (4.2)$$

- $\tilde{\alpha} = \alpha/h^2$  is the time-step-scaled compliance.
- XPBD solves this coupled system by a Gauss–Seidel sweep over constraints: for each constraint  $j$ , the multiplier increment is

$$\Delta \lambda_j = \frac{-C_j(\mathbf{x}) - \tilde{\alpha}_j \lambda_j}{\nabla C_j^\top \mathbf{M}^{-1} \nabla C_j + \tilde{\alpha}_j}, \quad (4.3)$$

- Corresponding position correction:  $\Delta \mathbf{x} = \mathbf{M}^{-1} \nabla C_j^\top \Delta \lambda_j$ .
- Velocities are recovered at substep end by  $\mathbf{v}_{t+h} = (\mathbf{x}_{t+h} - \mathbf{x}_t)/h$ .

#### 4.1.2 Algorithm

```

def xpbdd_substep(x, v, h, constraints):
    y = x + h * v + h**2 * M_inv @ f_ext
    x_new = y.copy()
    lambdas = zeros(len(constraints))
    for it in range(n_iter):
        for j, C_j in enumerate(constraints):
            dlam = (-C_j.value(x_new) - alpha_tilde[j] * lambdas[j]) / \
                (C_j.grad(x_new) @ M_inv @ C_j.grad(x_new).T
                 + alpha_tilde[j])
            x_new += M_inv @ C_j.grad(x_new).T * dlam
            lambdas[j] += dlam
    v_new = (x_new - x) / h
    return x_new, v_new

```

**Listing 4.1:** XPBD substep with  $n_{\text{iter}}$  Gauss–Seidel iterations.

### 4.1.3 Characteristics

- *Dual-variable.* An iteration’s state is encoded in the multipliers  $\lambda$ ; positions are derived.
- *Stiffness-by-compliance.* Material stiffness enters only through  $\tilde{\alpha} = 1/(kh^2)$ , so very stiff constraints naturally degenerate to projection ( $\tilde{\alpha} \rightarrow 0$ ).
- *First-order gradient information only.* Each constraint contributes its gradient  $\nabla C_j$ ; no Hessian of  $E$  is built.
- *Substeps preferred over iterations.* Following the small-steps result, I use many short substeps with one inner iteration.

## 4.2 VBD: primal block-coordinate descent

### 4.2.1 Derivation

- VBD performs block-coordinate descent on  $G$  in (2.5) with one block per vertex.
- Holding all other vertices fixed, the update for vertex  $i$  is the Newton step on the local 3-vector problem:

$$\mathbf{x}_i \leftarrow \mathbf{x}_i - \mathbf{H}_i^{-1} \mathbf{g}_i, \quad (4.4)$$

where

$$\mathbf{g}_i = \frac{m_i}{h^2} (\mathbf{x}_i - \mathbf{y}_i) + \sum_{e \in \mathcal{N}(i)} \nabla_i E_e(\mathbf{x}), \quad (4.5)$$

$$\mathbf{H}_i = \frac{m_i}{h^2} \mathbf{I}_3 + \sum_{e \in \mathcal{N}(i)} \nabla_{ii}^2 E_e(\mathbf{x}). \quad (4.6)$$

- For a distance-spring energy  $E_e = \frac{1}{2}k (\|\mathbf{x}_i - \mathbf{x}_j\| - \ell_0)^2$ , the per-vertex gradient and Hessian are closed-form and inexpensive.

### 4.2.2 Algorithm

```

def vbd_step(x, v, h):
    y = x + h * v + h**2 * M_inv @ f_ext
    x_new = y.copy() # inertial initialisation
    for it in range(n_iter):
        for i in vertex_order:
            g_i = (m[i] / h**2) * (x_new[i] - y[i])
            H_i = (m[i] / h**2) * I3
            for e in incident_edges(i):
                g_i += spring_grad_i(x_new, e)
                H_i += spring_hess_ii(x_new, e)
            x_new[i] -= solve_3x3(H_i, g_i)
    v_new = (x_new - x) / h
    return x_new, v_new

```

**Listing 4.2:** VBD step with  $n_{\text{iter}}$  block iterations.

### 4.2.3 Characteristics

- *Primal.* An iteration's state is positions  $\mathbf{x}$ ; no auxiliary multipliers.
- *Second-order locally.* Each block update is a Newton step on the true elastic Hessian restricted to one vertex.
- *Iterations over substeps.* VBD is normally used with one long step subdivided by many block iterations; the implicit Hessian handles stiffness without tiny  $h$ .
- *Parallel-friendly.* Vertices sharing no constraint can be updated independently (graph colouring). I disable parallelism for the core comparison; see Section 5.6.

## 4.3 Side-by-side: where the algorithms differ

|                       | XPBD                          | VBD                       |
|-----------------------|-------------------------------|---------------------------|
| Variable updated      | multipliers $\lambda$         | positions $\mathbf{x}$    |
| Order of information  | gradient $\nabla C$           | gradient + Hessian of $E$ |
| Block size            | per-constraint                | per-vertex                |
| Stiffness handled by  | compliance $\tilde{\alpha}$   | implicit Hessian          |
| Typical configuration | many substeps, 1 iteration    | one step, many iterations |
| Parallelism strategy  | constraint colouring (harder) | vertex colouring (easier) |

**Table 4.1:** Algorithmic contrast between XPBD and VBD.

- The row of Table 4.1 that drives the rest of the thesis is the first: XPBD lives in the dual, VBD lives in the primal.
- Every hypothesis in Section 1.3 ultimately traces back to that row.

## Chapter 5

# Methodology and implementation

- This chapter describes the implementation that produces the data in Chapter 6.
- More importantly, it describes the decisions that keep the XPBD-vs-VBD comparison fair.

### 5.1 Platform and environment

- Simulators implemented inside the Unity Editor.
- All measurements in Chapter 6 were collected in this environment:
  - CPU: *AMD Ryzen 7 5800H*.
  - RAM: *16 GB DDR4*.
  - GPU: *NVIDIA GeForce RTX 3060 Laptop GPU*.
  - OS: *Windows 11 Pro*.
  - Dev Tool: *Unity 6000.4.0f1*.

TODO: used Unity functionalities: Recorder, shader graph,

### 5.2 Integration architecture

- Both solvers share a common scene scaffolding.
- Unity drives a fixed-rate physics tick at  $1/dt = 30$  Hz; within each tick the active solver runs its update loop — substeps and inner iterations for XPBD, block iterations for VBD — on the shared particle/spring data structures.

The scaffolding handles:

- mesh and spring topology construction at scene load,
- particle/spring storage in flat C# arrays to keep cache behaviour comparable,
- recording of per-frame timing, energies, and reference solutions,
- rendering, which runs on the Unity main thread separately from the simulator.

### Overhead of the shared scaffolding.

- The scaffolding adds a small constant per-frame overhead (buffer swapping, Unity message dispatch, rendering).
- Both solvers pay this overhead identically, so it does not bias the comparison.
- I report the isolated solver time (excluding render and dispatch) separately from end-to-end frame time in Section 6.1.

## 5.3 How the comparison is kept fair

- A common pitfall is to set some parameter “the same” that is not actually the same physical quantity on the two sides.
- I adopt four explicit fairness controls.

### 5.3.1 Stiffness $\leftrightarrow$ compliance mapping

- VBD takes a stiffness  $k$  directly.
- XPBD takes a compliance  $\alpha = 1/k$ , scaled by the time step to  $\tilde{\alpha} = \alpha/h^2$ .
- To match materials I set:

$$\alpha_{\text{XPBD}} = 1/k_{\text{VBD}}, \tag{5.1}$$

- Mapping verified in a controlled test: a single spring with known analytic period simulated under both solvers; measured period required to match to within *TODO: tolerance, e.g. 1%*.
- Verification result reported in Section 6.3.

### 5.3.2 Matched compute budget — and the substep/iteration caveat

- The natural unit of work differs: XPBD’s primary lever is substep count, VBD’s is inner iteration count.
- **A substep is not an iteration.** An XPBD substep recomputes the inertial target  $\mathbf{y}$  and solves a fresh constraint system; a VBD iteration refines the same step.
- Because of this asymmetry, *wall-clock* is the primary axis of comparison throughout Chapter 6.
- I sweep substep counts for XPBD and iteration counts for VBD independently, plot quality vs. ms/frame, and read the *equal-wall-clock frontier* from the resulting curves.
- Iteration-count comparisons are reported only as secondary information.

### 5.3.3 Matched gravity, $dt$ , and damping

For all experiments unless otherwise noted:

- gravity  $\mathbf{g} = (0, -9.81, 0)$  m/s<sup>2</sup>,
- simulation time step  $dt = 1/30$  s,
- Rayleigh damping with matched  $\alpha_{\text{Rayleigh}}$  and  $\beta_{\text{Rayleigh}}$  coefficients (see Section 5.4),
- particle mass  $m_i = 1$  kg unless a mass-ratio sweep is in progress.

### 5.3.4 Disabling VBD’s accelerations for the core comparison

- The 2024 VBD paper [1] reports significant speedups from Chebyshev acceleration and adaptive initialisation.
- These features are genuinely part of VBD, but they confound an apples-to-apples comparison: an XPBD with Chebyshev-style acceleration would be an entirely different solver (essentially Anderson-accelerated XPBD).
- For the core comparison I therefore **disable Chebyshev acceleration and adaptive initialisation in VBD**, using only inertial initialisation.
- Results *with* the accelerations are reported separately in ??, framed as “what VBD looks like in production” rather than “what VBD looks like in this controlled comparison”.

## 5.4 Damping models

- Both solvers implement Rayleigh damping, but through different mechanisms.

### XPBD damping ( $\gamma$ -formulation).

- I use the multiplier-update extension of [6]:

$$\Delta\lambda_j = \frac{-C_j - \tilde{\alpha}_j\lambda_j - \tilde{\gamma}_j \nabla C_j^\top (\mathbf{x} - \mathbf{x}_{\text{prev}})}{\nabla C_j^\top \mathbf{M}^{-1} \nabla C_j (1 + \tilde{\gamma}_j) + \tilde{\alpha}_j}, \quad (5.2)$$

- $\tilde{\gamma}$  encodes the dissipation coefficient.
- This couples damping into the dual update without changing the algorithm’s structure.

### VBD damping (implicit Hessian).

- VBD absorbs Rayleigh damping into the local Hessian  $\mathbf{H}_i \rightarrow \mathbf{H}_i + \beta \mathbf{H}_{i,\text{el}}$  and into the gradient via a velocity term.
- Because the local solve is a Newton step, damping is handled fully implicitly with no additional outer iteration.
- Whether the two damping models produce the same physical decay at matched coefficients is analysed in Section 6.3.5, where it doubles as a fairness check.



## 5.5 Self-collision handling

The two solvers implement self-collision differently, by necessity:

- *XPBD self-collision* uses non-penetration constraints projected each iteration — contact enforced as a position constraint.
- *VBD self-collision* uses an elastic penalty potential added to  $E$ , contributing a gradient and Hessian to the affected vertex blocks.
- This is a model difference, not a solver difference: a penalty model *could* be implemented in XPBD and a projection model *could* be implemented in VBD, but the natural form differs.
- Flagged here so the self-collision result in Section 6.2.4 is read correctly: differences there are partly attributable to the contact model, not solely to XPBD vs. VBD.

## 5.6 Parallelisation (forward-looking)

- All measurements in this thesis are single-threaded on CPU.
- Noted here, and revisited in Section 7.3, the two solvers have different parallelisation profiles:
  - *VBD* parallelises naturally over vertex blocks via graph colouring: vertices sharing no constraint can be updated independently within a colour class. Well-suited to GPU execution.
  - *XPBD* can be parallelised via constraint-graph colouring, but the colouring is finer-grained (per constraint, not per vertex) and the dependency structure is less regular.
- The expected result — parallelisation favours VBD more than XPBD — is plausible but unverified by my experiments, and stands as a threat to external validity for any claim about which solver is “faster overall”.

# Chapter 6

## Experiments and results

- This chapter reports measurements along the four evaluation axes from Section 1.4: efficiency, robustness, physical realism, flexibility.
- All experiments use the shared scaffolding and fairness controls of Section 5.3: matched gravity,  $dt = 1/30$  s, no damping,  $\alpha_{\text{XPBD}} = 1/k_{\text{VBD}}$ , Chebyshev and adaptive initialisation disabled in VBD for the core comparison.

### 6.1 Computational efficiency

#### 6.1.1 Wall-clock

- Sweep through different substep and iteration configurations for each solver and plot their Pareto line (optimal solutions) based on avg ms/frame and avg??? Root Mean Squared position error from the Newton reference. Each diagram compares to a Newton with different substep counts
- XPBD substep count  $n_{\text{sub}} \in \{1, 2, 4, 8, 16, 32\}$
- VBD iteration count  $n_{\text{it}} \in \{1, 2, 4, 8, 16, 32\}$
- The simulated problem: a chain with 10 nodes, one end fixed, the other having a weight of ??? (with bending???) . Both ends start from the same height, only gravity acting on it.
- VBD reaches the same error value as XPBD with less iterations
- VBD's iterations cost more than XPBD's
- With this particular problem VBD's first guesses always seem to be closer to reference than XPBD's.
- From Newton with 4 substeps graph: VBD with the reference Newton's substeps converges really good. XPBD reaches it's smallest error with relatively few iterations, turning it higher doesn't improve error, only increases compute time.
- From Newton with more substeps: iterations loose their value as with small time steps both solvers converge at 1 or 2 iterations. XPBD takes over because of its cheaper iterations.

*TODO: add ss1, ss2, ss4, ss64*

### 6.1.2 Frame time vs. node count

- Measure FPS over a hanging-cloth scene with wind interaction at mesh resolutions *TODO: list resolutions, e.g.  $8 \times 8$  to  $128 \times 128$* , holding solver settings fixed at the Section 5.3 defaults.
- Based on the diagrams in ssec:ErrorVsMs, I choose substep 4 iteration 4 for XPBD and substep 4 iteration 2 for VBD, as these sit at the middle point of "realistic vs. efficient" and are always close to each other

*TODO: Add figure + table*

## 6.2 Robustness and stability

### 6.2.1 High mass ratio at the end of a pendulum

- A single spring pendulum with a heavy weight at the free end and a light chain leading to the pivot.
- Sweep the mass ratio  $m_{\text{tip}}/m_{\text{chain}}$  from 1 to  $10^4$ ; report each solver's qualitative behaviour and quantitative position error vs. Newton.
- *Hypothesis H1 prediction: XPBD degrades.*

*TODO: Figure — error vs. mass ratio.*

### 6.2.2 High stiffness ratio chain

- A chain with alternating stiff and soft springs; sweep the stiffness ratio from 1 to  $10^4$ .
- *Hypothesis H2 prediction: VBD degrades.*

*TODO: Figure — error vs. stiffness ratio.*

### 6.2.3 Self-collision: cloth folded onto itself

- A cloth dropped onto a small obstacle, folding onto itself.
- Qualitative visual comparison, noting that the contact model differs (projection vs. penalty) per Section 5.5.

*TODO: Renders, qualitative description.*

## 6.3 Physical realism

### 6.3.1 Newton reference

- Build a high-iteration Newton solver targeting  $\|\nabla G\| < 10^{-9}$  (or until further iterations reduce  $G$  by less than  $10^{-12}$ ).

- The Newton solver shares storage and energies with the production solvers, so the comparison is over algorithm, not material.
- Error metric throughout: mass-weighted position error  $\|\mathbf{x} - \mathbf{x}^*\|_{\mathbf{M}}$ , where  $\mathbf{x}^*$  is the Newton solution.

### 6.3.2 Harmonic oscillator

- A single spring-mass system oscillating in 1D, compared against the analytic solution  $x(t) = A \cos(\omega t - \phi) e^{-\zeta \omega t}$ .
- I measure:
  - period as a function of  $h$  and of iteration / substep count,
  - amplitude decay rate as a function of damping coefficient.
- This is the analogue of Fig. 2 in [6].

*TODO: Figure — measured period and decay, both solvers, against analytic.*

### 6.3.3 Energy tracking

- For an undamped free-falling and bouncing cloth, plot total mechanical energy over time.
- Identify systematic energy gain (instability) or loss (numerical damping) for each solver.

*TODO: Figure — energy vs. time, both solvers.*

### 6.3.4 Hanging cloth: stretched area vs. reference

- A square cloth pinned at the top corners, allowed to settle.
- Measure steady-state stretched area against the Newton reference.
- A realism check at static equilibrium, where transient dynamics drop out.

*TODO: Figure — stretched area error vs. stiffness.*

### 6.3.5 Damping cross-check

- At matched Rayleigh coefficients  $\alpha_{\text{Rayleigh}}$  and  $\beta_{\text{Rayleigh}}$ , measure the amplitude-decay curve produced by each solver on the harmonic oscillator.
- Identical decay confirms the damping models of Section 5.4 are calibrated consistently; this both validates fairness and is a result in its own right.

*TODO: Figure — decay envelopes overlaid.*

## 6.4 Flexibility

### 6.4.1 Iteration-count independence of stiffness

- For each solver, sweep stiffness  $k$  over several decades and measure how many iterations / substeps are needed to reach a fixed error threshold.
- Following Fig. 6 in [6], XPBD is expected to be roughly stiffness-independent in iteration count thanks to the compliance formulation.
- For both methods, stiffness stays roughly the same accross all iteration counts.

*TODO: Figure — iterations to threshold vs. stiffness.*

### 6.4.2 Stiffness-ratio test

- Same kind of plot as the previous subsection, but varying the *ratio* of stiffnesses across the mesh rather than a uniform stiffness.
- This is the experiment of [1] Fig. 24 (and AVBD [2] Fig. 4).
- Includes a brief discussion of AVBD's improvement on this axis without implementing it myself.

*TODO: Figure — error vs. stiffness ratio, with AVBD reference qualitatively cited.*

### 6.4.3 Ease of adding a new constraint or force

I implement a representative additional force — *TODO: choose, e.g. bending-angle constraint or area-preservation force* — in both solvers and report:

- lines of code added,
- whether a Hessian derivation was required (VBD: yes; XPBD: no, a gradient suffices),
- time to first correct simulation.
- The metric is admittedly subjective but captures a real practitioner cost that the more quantitative experiments above do not.

*TODO: Table — LOC, derivation effort, time-to-correct.*

# Chapter 7

## Discussion

- This chapter synthesises results across the four evaluation axes.
- It judges the hypotheses from Section 1.3 and enumerates the threats to validity that bound the conclusions.

### 7.1 Synthesis across axes

*TODO: once experimental results are in, write a paragraph per axis that explains how the result on that axis is consistent (or not) with the primal/dual framing of Section 2.5.*

- Efficiency: which solver dominates the equal-wall-clock frontier, and in what regime?
- Robustness: where do the stability envelopes diverge, and does the divergence match the primal/dual prediction?
- Realism: which solver’s converged state is closer to Newton at matched compute?
- Flexibility: how does the ease-of-extension cost compare against the per-frame cost benefit?

### 7.2 Verdicts on H1–H3

**H1 (XPBD degrades under high mass ratios) — *TODO: confirmed / refuted.***

- Tied to Section 6.2.1 and Section 6.2.2.
- Mechanism (if confirmed): dual updates propagate slowly through chains of strongly mismatched masses, because each Gauss–Seidel sweep changes only one constraint’s multiplier, and the change couples to the next constraint only through the position update on the shared particle.

**H2 (VBD degrades under high stiffness ratios) — *TODO: confirmed / refuted.***

- Tied to Section 6.2.3 and Section 6.4.2.
- Mechanism (if confirmed): the local Hessian block  $\mathbf{H}_i$  at a vertex sums contributions from every incident spring, so widely varying stiffnesses explode the block’s condition number, slowing the Newton step.

**H3 (VBD’s per-iteration convergence is more stiffness-robust) — *TODO: confirmed / refuted.***

- Tied to ?? and Section 6.4.1.
- Mechanism (if confirmed): VBD uses second-order information about  $E$  locally, so increasing  $k$  does not change the number of Newton steps to convergence much; XPBD discards this information.

**Joint reading.**

- If H1, H2, H3 all hold, the framing is consistent: each solver is robust exactly in the regime where its update direction sees the right information, and brittle where the other side of the dual would.
- The practical recommendation in that case is regime-based, not absolute — itself a useful contribution if the experiments support it.

## **7.3 Threats to validity**

### **7.3.1 Single-threaded execution**

- All measurements run single-threaded on CPU.
- VBD is known to parallelise better than XPBD (Section 5.6), so ms/frame results underestimate VBD’s likely production advantage.
- This affects external validity of any quantitative efficiency claim but does not invalidate qualitative claims about convergence quality or stability envelopes.

### **7.3.2 Mass-spring discretisation only**

- I use distance-spring energies throughout, including for bending; real cloth simulators use FEM-based stretch energies and dihedral bending.
- The primal/dual framing carries over (both XPBD and VBD admit arbitrary elastic energies), but the specific stiffness ratios encountered — and therefore the magnitude of the H2/H3 effects — may differ.

### **7.3.3 Gauss–Seidel order dependence**

- Both solvers are order-sensitive (Section 2.4).
- I fix a single deterministic order; an adversarial reordering could in principle shift results.
- I do not investigate this, and it is a real threat to reproducibility on other implementations.

### 7.3.4 Mismatched collision models

- XPBD uses projection-based self-collision; VBD uses penalty-based self-collision (Section 5.5).
- Any difference in the self-collision experiments mixes algorithmic and contact-model effects.

### 7.3.5 Fixed $dt = 1/30$ s

- $dt = 1/30$  s is generous for Unity but stresses both solvers.
- A different fixed  $dt$ , or an adaptive step, could shift the equal-wall-clock frontier.
- The stability map in ?? mitigates this by characterising both solvers over a  $dt$  range, but the headline results in Section 6.1 are tied to a single  $dt$ .

## 7.4 Limitations

Beyond the threats above, three further limitations are worth naming:

- I do not implement AVBD or any post-2024 VBD extension.
- Self-collision uses a simple model; production simulators often use continuous collision detection (CCD), which both solvers can be augmented with at additional cost.
- Friction is not modelled; the experiments are friction-free.



## Chapter 8

# Conclusion and future work

### 8.1 Summary

- This thesis presented a controlled comparison of XPBD and VBD, the two dominant solvers for real-time deformable simulation.
- The contribution is threefold:
  - a shared Unity codebase implementing both solvers under matched material parameters and compute budget;
  - a Newton ground-truth reference, enabling reporting of solution error rather than only relative performance;
  - an experimental design driven by the primal/dual theory of [5], organised around three falsifiable hypotheses (H1–H3).

*TODO: once results are finalised, a single paragraph summary of Section 7.2 goes here: which hypotheses hold, what the equal-wall-clock frontier looks like, and the regime-based recommendation for practitioners.*

### 8.2 Future work

#### 8.2.1 Multithreaded and GPU comparison

- The most consequential extension is to repeat the experiments on a multithreaded CPU and on the GPU.
- Both solvers admit parallelisation via graph colouring, but the structures differ in granularity (Section 5.6).
- Expected outcome: VBD’s per-vertex blocks colour more cleanly than XPBD’s per-constraint blocks, amplifying its production advantage; quantifying this gap is the natural next study.

#### 8.2.2 Extend research to other dynamic types

- Beside mass-spring systems, in theory both solvers could also handle other cloth representations, soft bodies, and rigid bodies.

- It would be interesting to see how the two solvers perform on these other types of dynamics, and whether the same conclusions hold.

### 8.2.3 FEM constraints

- Replacing mass-spring with FEM-based stretch and bending energies would lift the realism ceiling and test whether the H1–H3 effects survive the change of discretisation.
- The primal/dual framing predicts they should, but only experiments can confirm.

### 8.2.4 AVBD and post-2024 variants

- Augmented VBD [2] reportedly handles wide stiffness ratios substantially better than canonical VBD.
- Re-running the H2 and stiffness-ratio experiments with AVBD would clarify whether the gap I observe is intrinsic to the primal approach or specific to the canonical 2024 implementation.

# Acknowledgements

Ez nem kötelező, akár törölhető is. Ha a szerző szükségét érzi, itt lehet köszönetet nyilvánítani azoknak, akik hozzájárultak munkájukkal ahhoz, hogy a hallgató a szakdolgozatban vagy diplomamunkában leírt feladatokat sikeresen elvégezze. A konzulensnek való köszönetnyilvánítás sem kötelező, a konzulensnek hivatalosan is dolga, hogy a hallgatót konzultálja.

# Bibliography

- [1] Anka He Chen, Ziheng Liu, Yin Yang, and Cem Yuksel. Vertex block descent. *ACM Transactions on Graphics (Proceedings of SIGGRAPH 2024)*, 43(4):116:1–116:16, 07 2024. ISSN 0730-0301. DOI: 10.1145/3658179. URL <https://doi.org/10.1145/3658179>.
- [2] Chris Giles, Elie Diaz, and Cem Yuksel. Augmented vertex block descent. *ACM Transactions on Graphics (Proceedings of SIGGRAPH 2025)*, 44(4), 08 2025. ISSN 0730-0301. DOI: 10.1145/3731195. URL <https://doi.org/10.1145/3731195>.
- [3] Minchen Li, Chenfanfu Jiang, Zhaofeng Luo, Wenxin Du, Chang Yu, Žiga Kovačič, and Tianyi Xie. *Physics-Based Simulation*. March 2026. URL <https://phys-sim-book.github.io/>.
- [4] Jun Ma, Nak-Jun Sung, Min-Hyung Choi, and Min Hong. Performance comparison of vertex block descent and position based dynamics algorithms using cloth simulation in unity. *Applied Sciences*, 14(23), 2024. ISSN 2076-3417. DOI: 10.3390/app142311072. URL <https://www.mdpi.com/2076-3417/14/23/11072>.
- [5] M. Macklin, K. Erleben, M. Müller, N. Chentanez, S. Jeschke, and T. Y. Kim. Primal/dual descent methods for dynamics. In *Proceedings of the ACM SIGGRAPH/Eurographics Symposium on Computer Animation*, SCA '20, Goslar, DEU, 2020. Eurographics Association. DOI: 10.1111/cgf.14104. URL <https://doi.org/10.1111/cgf.14104>.
- [6] Miles Macklin, Matthias Müller, and Nuttapong Chentanez. Xpbd: position-based simulation of compliant constrained dynamics. In *Proceedings of the 9th International Conference on Motion in Games*, MIG '16, page 49–54, New York, NY, USA, 2016. Association for Computing Machinery. ISBN 9781450345927. DOI: 10.1145/2994258.2994272. URL <https://doi.org/10.1145/2994258.2994272>.
- [7] Miles Macklin, Kier Storey, Michelle Lu, Pierre Terdiman, Nuttapong Chentanez, Stefan Jeschke, and Matthias Müller. Small steps in physics simulation. *Proceedings of the 18th annual ACM SIGGRAPH/Eurographics Symposium on Computer Animation*, 2019. URL <https://api.semanticscholar.org/CorpusID:197928472>.
- [8] Matthias Müller. Ten minute physics. <https://matthias-research.github.io/pages/tenMinutePhysics/index.html>, 2022.
- [9] Matthias Müller, Bruno Heidelberger, Marcus Hennix, and John Ratcliff. Position based dynamics. *Journal of Visual Communication and Image Representation*, 18(2):109–118, 2007. ISSN 1047-3203. DOI: <https://doi.org/10.1016/j.jvcir.2007.01.005>. URL <https://www.sciencedirect.com/science/article/pii/S1047320307000065>.

d

# Appendix

## A.1 Declaration on the Use of Generative Artificial Intelligence

☐ **I have not used** any generative AI tools.

☐ **I have used generative AI tools.** I have verified the content generated by AI, ensured the accuracy of the outputs, and properly indicated each instance of use in the table below.

| Usage type                                                          | Name of Generative AI Tool(s) | Affected Sections<br>(chapter, page number, reference) | Estimate<br>Proportion of Use<br>(per usage type) |
|---------------------------------------------------------------------|-------------------------------|--------------------------------------------------------|---------------------------------------------------|
| Literature Review                                                   |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Program Code Generation                                             |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Generating New Ideas or Solution Proposals                          |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Creating an Outline<br>(text structure, bullet points)              |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Creating Text Blocks                                                |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Generating Images for Illustrative Purposes                         |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Data Visualization,<br>Generating Charts Based on Data Points       |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Preparing a Presentation                                            |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| Other (please specify)                                              |                               |                                                        |                                                   |
| <b>Brief Summary of the Prompt</b>                                  |                               |                                                        |                                                   |
| <b>Aggregated Percentage Value (for the core part of the task):</b> |                               |                                                        |                                                   |
| <b>Brief Textual Justification of the Aggregated Value:</b>         |                               |                                                        |                                                   |
|                                                                     |                               |                                                        |                                                   |